

Amazon_reviews_sentiment_analysis

August 18, 2021

```
[6]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")

import re
from bs4 import BeautifulSoup
from tqdm import tqdm
from nltk.stem import WordNetLemmatizer

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix

from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.layers import Dense, Input, LSTM, Embedding, Dropout, \
    ↪Activation, GRU, Flatten
from keras.layers import Bidirectional, GlobalMaxPool1D
from keras.models import Model, Sequential
from keras.layers import Convolution1D
from keras import initializers, regularizers, constraints, optimizers, layers
```

```
[7]: import io
from google.colab import files
uploaded = files.upload()
df2 = pd.read_csv(io.BytesIO(uploaded['amazon_reviews_1.csv']))
df = pd.read_csv('amazon_reviews_1.csv')
df.head()
```

<IPython.core.display.HTML object>

Saving amazon_reviews_1.csv to amazon_reviews_1 (1).csv

```
[7]: Unnamed: 0 ... review_category
0      35311 ...                1
1      44590 ...                0
```

```

2      90295 ...      0
3      38798 ...      0
4      43364 ...      0

```

[5 rows x 16 columns]

```
[8]: df.columns
```

```
[8]: Index(['Unnamed: 0', 'marketplace', 'customer_id', 'review_id', 'product_id',
          'product_parent', 'product_title', 'product_category', 'star_rating',
          'helpful_votes', 'total_votes', 'vine', 'verified_purchase',
          'review_date', 'review', 'review_category'],
          dtype='object')
```

```
[9]: l=['marketplace','product_category','verified_purchase','vine','review_category','helpful_votes',
      ↪ 'total_votes']
```

```

for i in l:
    print('The unique element in : ', i)
    print('-'*30)
    print(df[i].unique())
    print()

```

The unique element in : marketplace

['UK']

The unique element in : product_category

['Mobile_Apps' 'Video DVD' 'Music' 'Books']

The unique element in : verified_purchase

['Y' 'N']

The unique element in : vine

['N' 'Y']

The unique element in : review_category

[1 0]

The unique element in : helpful_votes

[0 1 3 5 14 2 4 29 32 9 6 13 20 19 17 7 22 15
18 28 31 8 10 11 190 100 30 16 42 12 38 35 102 95 27 69
37 49 36 84 133 68 184 62 43 26 24 44 21 419 33 77 25 109]

```

97 39 76 56 175 46 23 80 166 34 48 63 292 41 86 127 58 70
53 79 45 229 163 51 380 105 47 297 244 55 117 103 50 75 135 101
90]

```

The unique element in : total_votes

```

-----
[ 1  2  5  0 33 41 10  3  4 19 11 45  8 31  6 84  7 15
 23 134  9 12 14 21 27 18 28 22 13 16 25 29 109 80 26 52
 17 36 128 20 216 113 104 35 30 37 48 34 24 42 38 96 69 50
 82 110 57 46 108 98 121 87 53 91 64 39 60 62 90 142 92 196
 79 94 77 66 59 493 43 89 49 40 56 85 70 219 83 88 44 202
 54 78 178 169 72 147 32 74 305 67 47 97 149 65 55 201 132 112
 71 268 51 189 434 118 58 335 125 290 63 61 75 163 123 68 151 103
130 222 81]

```

```

[10]: l=['product_category','verified_purchase','vine','review_category','helpful_votes',
        ↪'total_votes']

for i in l:
    print('The unique element in : ', i)
    print('-'*30)
    print(df[i].value_counts())
    print()

```

The unique element in : product_category

```

-----
Books          2624
Music          2493
Video DVD      2469
Mobile_Apps    2414
Name: product_category, dtype: int64

```

The unique element in : verified_purchase

```

-----
Y      8119
N      1881
Name: verified_purchase, dtype: int64

```

The unique element in : vine

```

-----
N      9990
Y        10
Name: vine, dtype: int64

```

The unique element in : review_category

```

-----
0      5046

```

```

1      4954
Name: review_category, dtype: int64

The unique element in : helpful_votes
-----
0      6891
1      1183
2       475
3       298
4       207
...
58         1
53         1
77         1
101        1
175        1
Name: helpful_votes, Length: 91, dtype: int64

The unique element in : total_votes
-----
0      5205
1      1411
2       765
3       431
4       314
...
202         1
178         1
130         1
82          1
123         1
Name: total_votes, Length: 129, dtype: int64

```

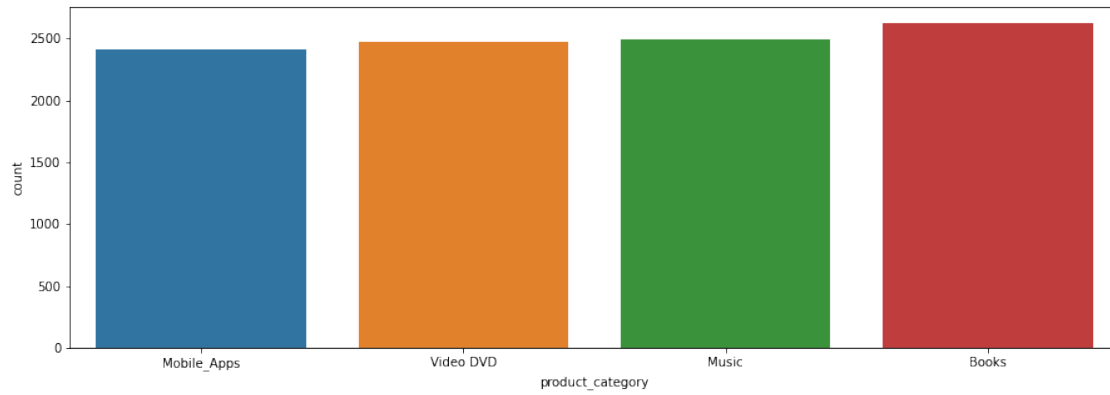
```

[11]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
l=['product_category','verified_purchase','vine','review_category','helpful_votes',
  ↪ 'total_votes']

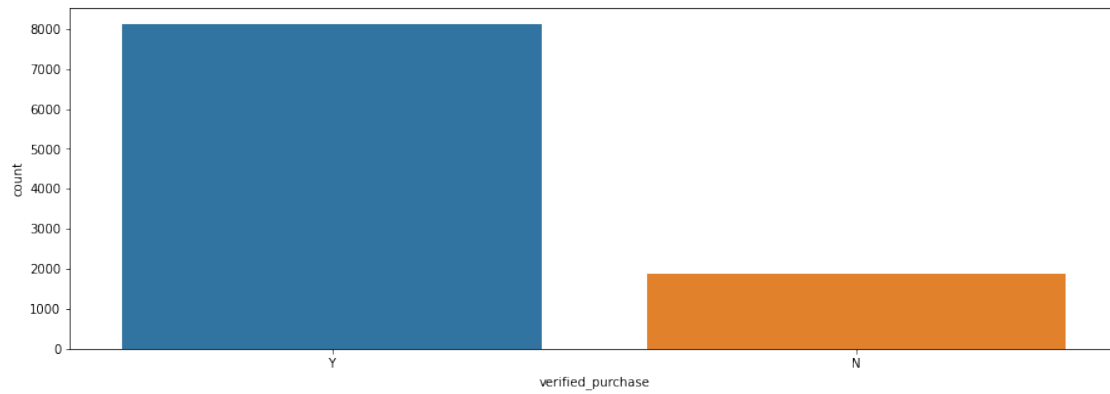
for i in l:
    print('The count plot of  : ', i)
    print('-'*30)
    plt.figure(figsize=(15,5))
    sns.countplot(df[i])
    plt.show()
    print()

```


The count plot of : product_category



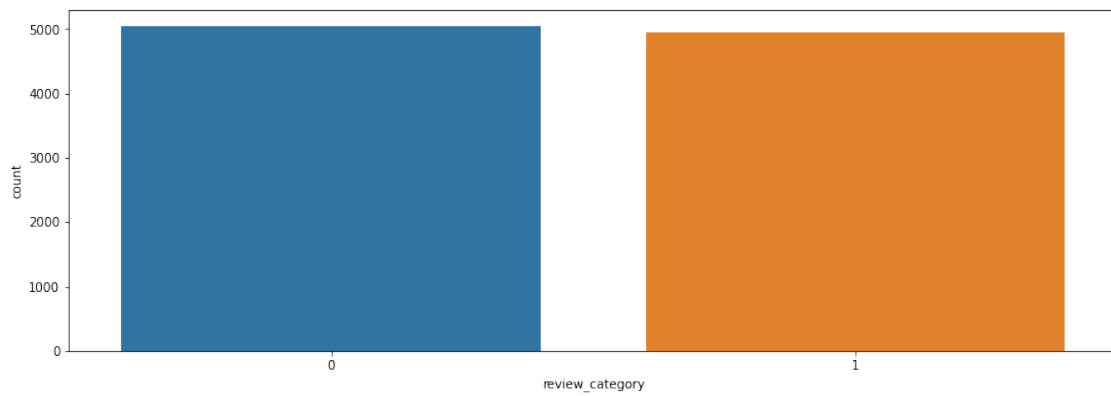
The count plot of : verified_purchase




The count plot of : vine




The count plot of : review_category



The count plot of : helpful_votes



The count plot of : total_votes



```
[12]: df_new=df[['star_rating','review']]
      df_new.head()
```

```
[12]:   star_rating      review
0         5  Addictive game Very good game, the graphics ar...
1         2  Where to start? I came on here because I wante...
2         3  Really Disappointed in this Film I have to say...
3         1  Rubbish game When I tried to open it, it said ...
4         1  One Star Come on
```

```
[13]: df_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   star_rating  10000 non-null  int64
1   review       9999 non-null   object
dtypes: int64(1), object(1)
memory usage: 156.4+ KB
```

```
[14]: df_new.dropna(inplace=True)
      df_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 9999 entries, 0 to 9999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   star_rating  9999 non-null   int64
1   review       9999 non-null   object
dtypes: int64(1), object(1)
memory usage: 234.4+ KB
```


```
[15]: df_new['star_rating'].value_counts()
```

```
[15]: 5    4149
      3    2350
      1    1661
      2    1035
      4     804
      Name: star_rating, dtype: int64
```

```
[16]: #show the visualizing count of target feature.
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

sns.countplot(df_new['star_rating'])
```

```
[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb70247e6d0>
```




```
[17]: df_new.head(1)
```

```
[17]:   star_rating      review
0           5  Addictive game Very good game, the graphics ar...
```


```
[18]: df_new["rating"] = df_new["star_rating"].apply(lambda x : 1 if x>3 else 0)
df_new.head(3)
```

```
[18]:   star_rating      review  rating
0           5  Addictive game Very good game, the graphics ar...      1
1           2  Where to start? I came on here because I want...      0
2           3  Really Disappointed in this Film I have to say...      0
```

```
[19]: #show the visualizing count of target feature.
```

```
sns.countplot(df_new['rating'])
```

```
[19]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb7023e2e90>
```



4. Text preprocessing In this step, following operations are performed on the review text

- Removing website links
- Removing html tags
- Decontracting(expanding from the original form)

- Removing the words with numeric digits
- Removing non-word characters
- Converting to lower case
- Removing stop words
- Performing Lemmatization

```
[20]: def decontract(text):
    text = re.sub(r"won't", "will not", text)
    text = re.sub(r"can't", "can not", text)
    text = re.sub(r"n't", " not", text)
    text = re.sub(r"\ 're", " are", text)
    text = re.sub(r"\ 's", " is", text)
    text = re.sub(r"\ 'd", " would", text)
    text = re.sub(r"\ 'll", " will", text)
    text = re.sub(r"\ 't", " not", text)
    text = re.sub(r"\ 've", " have", text)
    text = re.sub(r"\ 'm", " am", text)
    return text
```

```
[21]: lemmatizer = WordNetLemmatizer()
import nltk
nltk.download('stopwords')
nltk.download('wordnet')
from nltk.corpus import stopwords
stopwords=set(stopwords.words('english'))
def preprocess_text(review):
    review = re.sub(r"http\S+", "", review) # removing website links
    review = BeautifulSoup(review, 'lxml').get_text() # removing html tags
    review = decontract(review) # decontracting
    review = re.sub("\S*\d\S*", "", review).strip() # removing the words
    ↪with numeric digits
    review = re.sub('[^A-Za-z]+', ' ', review) # removing non-word
    ↪characters
    review = review.lower() # converting to lower
    ↪case
    review = [word for word in review.split(" ") if not word in stopwords] #
    ↪removing stop words
    review = [lemmatizer.lemmatize(token, "v") for token in review]
    ↪#Lemmatization
    review = " ".join(review)
    review.strip()
    return review
df_new['cleaned_text'] = df_new['review'].apply(lambda x: preprocess_text(x))
```

[nltk_data] Downloading package stopwords to /root/nltk_data...

[nltk_data] Unzipping corpora/stopwords.zip.

[nltk_data] Downloading package wordnet to /root/nltk_data...

[nltk_data] Unzipping corpora/wordnet.zip.

```
[22]: df_new.head(20)
```

```
[22]:   star_rating  ... cleaned_text
0          5  ... addictive game good game graphics brilliant gr...
1          2  ... start come want get flappy bird saw think basi...
2          3  ... really disappoint film say look forward remake...
3          1  ...      rubbish game try open say stop fix higher rat
4          1  ...                                one star come
5          5  ...      five star brilliant fast postage fab thank
6          5  ... five star cd excellent wish review system woul...
7          1  ... worth buy family huge matt damon fan especiall...
8          5  ...                                five star good buy
9          1  ...                                one star like
10         1  ... hand hold shakey camera rubbish big monster at...
11         5  ... fantastic film great especially original end d...
12         5  ...                                love love game
13         5  ...      five star great book recommend anyone
14         4  ...      four star really make difference try complete
15         5  ...      five star beautiful cd would recomened
16         3  ...                                worth alook good hop
17         1  ... disappointingly flat album take hostile muse a...
18         3  ... bone booth happy families bone booth settle ha...
19         5  ... best booko ever best book ever love read chang...
```

[20 rows x 4 columns]


```
[23]: df_final=df_new[['rating','cleaned_text']]
      df_final.head()
```


```
[23]:   rating  ... cleaned_text
0        1  ... addictive game good game graphics brilliant gr...
1        0  ... start come want get flappy bird saw think basi...
2        0  ... really disappoint film say look forward remake...
3        0  ...      rubbish game try open say stop fix higher rat
4        0  ...                                one star come
```

```
[24]: #most common words used
      from wordcloud import WordCloud


      for i in range (0,2):
          print('When star Rating is : ', i)
          print('_'*30)
          print()
          text = df_final[df_final['rating'] == i]
          all_words = ' '.join([text for text in text.cleaned_text])
          wordcloud = WordCloud(width= 1500, height= 800,
                                max_font_size = 170,
```

When star Rating is : 0





When star Rating is : 1



```
[26]: train_df, test_df = train_test_split(df_final, test_size = 0.2, random_state = 42)
print("Training data size : ", train_df.shape)
print("Test data size : ", test_df.shape)
```

Training data size : (7999, 2)

Test data size : (2000, 2)

6. Model Building 6.1 Fitting LSTM with Embedding layer

```
[27]: top_words = 6000
tokenizer = Tokenizer(num_words=top_words)
tokenizer.fit_on_texts(train_df['cleaned_text'])
list_tokenized_train = tokenizer.texts_to_sequences(train_df['cleaned_text'])

max_review_length = 130
X_train = pad_sequences(list_tokenized_train, maxlen=max_review_length)
y_train = train_df['rating']
```

```
[28]: embedding_vecor_length = 32
model = Sequential()
model.add(Embedding(top_words+1, embedding_vecor_length,
    input_length=max_review_length))
model.add(LSTM(200))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam',
    metrics=['accuracy'])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 130, 32)	192032
lstm (LSTM)	(None, 200)	186400
dense (Dense)	(None, 1)	201

Total params: 378,633
 Trainable params: 378,633
 Non-trainable params: 0

```
[29]: model.fit(X_train,y_train, epochs=20, batch_size=8, validation_split=0.2)
```

Epoch 1/20

800/800 [=====] - 30s 12ms/step - loss: 0.5203 - accuracy: 0.7483 - val_loss: 0.2719 - val_accuracy: 0.8900

Epoch 2/20
800/800 [=====] - 8s 10ms/step - loss: 0.2259 -
accuracy: 0.9153 - val_loss: 0.2742 - val_accuracy: 0.8894
Epoch 3/20
800/800 [=====] - 8s 10ms/step - loss: 0.1381 -
accuracy: 0.9510 - val_loss: 0.3179 - val_accuracy: 0.8806
Epoch 4/20
800/800 [=====] - 8s 10ms/step - loss: 0.0926 -
accuracy: 0.9682 - val_loss: 0.3316 - val_accuracy: 0.8850
Epoch 5/20
800/800 [=====] - 8s 11ms/step - loss: 0.0615 -
accuracy: 0.9797 - val_loss: 0.3978 - val_accuracy: 0.8844
Epoch 6/20
800/800 [=====] - 8s 10ms/step - loss: 0.0546 -
accuracy: 0.9843 - val_loss: 0.4403 - val_accuracy: 0.8813
Epoch 7/20
800/800 [=====] - 8s 10ms/step - loss: 0.0662 -
accuracy: 0.9794 - val_loss: 0.4784 - val_accuracy: 0.8719
Epoch 8/20
800/800 [=====] - 8s 10ms/step - loss: 0.0332 -
accuracy: 0.9906 - val_loss: 0.5904 - val_accuracy: 0.8756
Epoch 9/20
800/800 [=====] - 8s 10ms/step - loss: 0.0230 -
accuracy: 0.9937 - val_loss: 0.5682 - val_accuracy: 0.8737
Epoch 10/20
800/800 [=====] - 8s 11ms/step - loss: 0.0176 -
accuracy: 0.9950 - val_loss: 0.6010 - val_accuracy: 0.8675
Epoch 11/20
800/800 [=====] - 8s 10ms/step - loss: 0.0248 -
accuracy: 0.9927 - val_loss: 0.6227 - val_accuracy: 0.8606
Epoch 12/20
800/800 [=====] - 8s 10ms/step - loss: 0.0176 -
accuracy: 0.9965 - val_loss: 0.7198 - val_accuracy: 0.8662
Epoch 13/20
800/800 [=====] - 8s 10ms/step - loss: 0.0124 -
accuracy: 0.9961 - val_loss: 0.5921 - val_accuracy: 0.8650
Epoch 14/20
800/800 [=====] - 8s 10ms/step - loss: 0.0152 -
accuracy: 0.9952 - val_loss: 0.7029 - val_accuracy: 0.8687
Epoch 15/20
800/800 [=====] - 8s 10ms/step - loss: 0.0156 -
accuracy: 0.9954 - val_loss: 0.6552 - val_accuracy: 0.8644
Epoch 16/20
800/800 [=====] - 8s 10ms/step - loss: 0.0111 -
accuracy: 0.9961 - val_loss: 0.8154 - val_accuracy: 0.8650
Epoch 17/20
800/800 [=====] - 8s 10ms/step - loss: 0.0043 -
accuracy: 0.9986 - val_loss: 0.7481 - val_accuracy: 0.8656


```
Epoch 18/20
800/800 [=====] - 8s 10ms/step - loss: 0.0089 -
accuracy: 0.9978 - val_loss: 0.8074 - val_accuracy: 0.8687
Epoch 19/20
800/800 [=====] - 8s 10ms/step - loss: 0.0065 -
accuracy: 0.9978 - val_loss: 0.7939 - val_accuracy: 0.8694
Epoch 20/20
800/800 [=====] - 8s 10ms/step - loss: 0.0055 -
accuracy: 0.9988 - val_loss: 0.8488 - val_accuracy: 0.8637
```

[29]: <keras.callbacks.History at 0x7fb6fdc73650>

0.1 Adding dropout

```
[30]: embedding_vecor_length = 32
model = Sequential()
model.add(Embedding(top_words+1, embedding_vecor_length,
    ↪input_length=max_review_length))
model.add(Dropout(0.3))
model.add(LSTM(100))
model.add(Dropout(0.3))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam',
    ↪metrics=['accuracy'])
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 130, 32)	192032
dropout (Dropout)	(None, 130, 32)	0
lstm_1 (LSTM)	(None, 100)	53200
dropout_1 (Dropout)	(None, 100)	0
dense_1 (Dense)	(None, 1)	101

=====
 Total params: 245,333
 Trainable params: 245,333
 Non-trainable params: 0
 =====

```
[31]: model.fit(X_train,y_train, epochs=5, batch_size=64, validation_split=0.2)
```

Epoch 1/5

```

100/100 [=====] - 4s 19ms/step - loss: 0.6247 -
accuracy: 0.6407 - val_loss: 0.3311 - val_accuracy: 0.8644
Epoch 2/5
100/100 [=====] - 1s 12ms/step - loss: 0.3325 -
accuracy: 0.8688 - val_loss: 0.2945 - val_accuracy: 0.8781
Epoch 3/5
100/100 [=====] - 1s 12ms/step - loss: 0.2581 -
accuracy: 0.8956 - val_loss: 0.2759 - val_accuracy: 0.8813
Epoch 4/5
100/100 [=====] - 1s 13ms/step - loss: 0.1736 -
accuracy: 0.9353 - val_loss: 0.2735 - val_accuracy: 0.8844
Epoch 5/5
100/100 [=====] - 1s 13ms/step - loss: 0.1403 -
accuracy: 0.9505 - val_loss: 0.2966 - val_accuracy: 0.8800

```

[31]: <keras.callbacks.History at 0x7fb6f0423e50>

6.2 Evaluating model performance on test data

```

[32]: list_tokenized_test = tokenizer.texts_to_sequences(test_df['cleaned_text'])
X_test = pad_sequences(list_tokenized_test, maxlen=max_review_length)
y_test = test_df['rating']
prediction = model.predict(X_test)
y_pred = (prediction > 0.5)
print("Accuracy of the model : ", accuracy_score(y_pred, y_test))
print('F1-score: ', f1_score(y_pred, y_test))
print('Confusion matrix:')
confusion_matrix(y_test,y_pred)

```

```

Accuracy of the model :  0.8765
F1-score:  0.8711528429838289
Confusion matrix:

```

[32]: array([[918, 111],
[136, 835]])

1 ML

```

[33]: df_final.head(2)

```

```

[33]:      rating      cleaned_text
0         1  addictive game good game graphics brilliant gr...
1         0  start come want get flappy bird saw think basi...

```

```

[34]: x=df_final['cleaned_text']
y=df_final['rating']
print(x[:5])
print(y[:5])

```

```

0    addictive game good game graphics brilliant gr...
1    start come want get flappy bird saw think basi...
2    really disappoint film say look forward remake...
3        rubbish game try open say stop fix higher rat
4                                one star come
Name: cleaned_text, dtype: object
0    1
1    0
2    0
3    0
4    0
Name: rating, dtype: int64

```

```

[35]: #countvectoriser with stemming

from sklearn.feature_extraction.text import CountVectorizer
CV=CountVectorizer(max_features=5000, ngram_range=(2,2))
x_stem=CV.fit_transform(x).toarray()
x_stem

```

```

[35]: array([[0, 0, 0, ..., 0, 0, 0],
            [0, 0, 0, ..., 0, 0, 0],
            [0, 0, 0, ..., 0, 0, 0],
            ...,
            [0, 0, 0, ..., 0, 0, 0],
            [0, 0, 0, ..., 0, 0, 0],
            [0, 0, 0, ..., 0, 0, 0]])

```

```

[36]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test= train_test_split(x_stem,y,test_size=0.2)

```

```

[37]: X_train.shape, X_test.shape, y_train.shape, y_test.shape

```

```

[37]: ((7999, 5000), (2000, 5000), (7999,), (2000,))

```

```

[38]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import SGDClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import ExtraTreesClassifier

from sklearn import metrics

```

```
[39]: clf_logit_pca = LogisticRegression()
      clf_tree_pca = DecisionTreeClassifier()
      clf_svm_pca = SVC(probability = True)
      clf_rf_pca= RandomForestClassifier(n_estimators=100)
      clf_gb_pca= GradientBoostingClassifier(n_estimators=100)
      clf_bag_pca= BaggingClassifier(n_estimators=100)
      clf_nb_pca= GaussianNB()
      clf_sgd_pca= SGDClassifier(loss='modified_huber',shuffle=True, random_state=101)
      clf_knn_pca= KNeighborsClassifier(n_neighbors=15)
      clf_etc_pca= ExtraTreesClassifier(n_estimators = 5,
                                       criterion = 'entropy', max_features = 2)
```

```
[40]: classifiers = ['LogisticRegression', 'DecisionTreeClassifier', 'SVC',
                    ↪ 'RandomForestClassifier',
                    'GradientBoostingClassifier', 'BaggingClassifier', 'GaussianNB',
                    ↪ 'SGDClassifier',
                    'KNeighborsClassifier', 'ExtraTreesClassifier' ]
```

```
[41]: models = {clf_logit_pca: 'LogisticRegression',
                 clf_tree_pca: 'DecisionTreeClassifier',
                 clf_svm_pca: 'SVC',
                 clf_rf_pca: 'RandomForestClassifier',
                 clf_gb_pca: 'GradientBoostingClassifier',
                 clf_bag_pca: 'BaggingClassifier',
                 clf_nb_pca: 'GaussianNaiveBayes',
                 clf_sgd_pca: 'StochasticGradientDescent',
                 clf_knn_pca: 'KNN',
                 clf_etc_pca: 'ExtraTreesClassifier'}
```

```
[42]: def train(algo, name, X_train, y_train, X_test, y_test):
      algo.fit(X_train, y_train)
      y_test_pred = algo.predict(X_test)
      return y_test_pred

      def performance(algo, name, X_test, y_test, y_test_pred):
          # Accuracy Score
          score = metrics.accuracy_score(y_test, y_test_pred)
          print('Accuracy score for', name, ':', score)

          # Confusion Metrics
          sns.heatmap(metrics.confusion_matrix(y_test, y_test_pred), annot=True,
          ↪fmt='d')
          plt.title("Confusion Matrix")
          plt.show()

          # Classification Report
```

```

print(metrics.classification_report(y_test, y_test_pred))

# ROC AUC Score
prob = algo.predict_proba(X_test)
prob = prob[:, 1]
roc_auc = metrics.roc_auc_score(y_test, prob)
print(roc_auc)

# Plotting ROC Curve
fpr, tpr, thresholds = metrics.roc_curve(y_test, prob)
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.show()

return roc_auc

def main(models):
    auc = []
    for algo, name in models.items():
        print('#'*40, name, '#'*40)
        y_test_pred = train(algo, name, X_train, y_train, X_test, y_test)
        roc_auc = performance(algo, name, X_test, y_test, y_test_pred)
        auc.append(roc_auc)
    return auc

```


[43]: %%time

```
auc = main(models)
```

```


##### LogisticRegression
#####
Accuracy score for LogisticRegression : 0.8265

```




	precision	recall	f1-score	support
0	0.80	0.88	0.83	997
1	0.86	0.78	0.82	1003
accuracy			0.83	2000
macro avg	0.83	0.83	0.83	2000
weighted avg	0.83	0.83	0.83	2000

0.9162422461802157




```
##### DecisionTreeClassifier
#####
Accuracy score for DecisionTreeClassifier : 0.7495
```




	precision	recall	f1-score	support
0	0.83	0.62	0.71	997
1	0.70	0.87	0.78	1003
accuracy			0.75	2000
macro avg	0.77	0.75	0.75	2000
weighted avg	0.77	0.75	0.75	2000

0.8037397336576029




```
##### SVC
#####
Accuracy score for SVC : 0.813
```





	precision	recall	f1-score	support
0	0.75	0.94	0.83	997
1	0.92	0.69	0.79	1003
accuracy			0.81	2000
macro avg	0.83	0.81	0.81	2000
weighted avg	0.83	0.81	0.81	2000

0.9135582220239982




```
##### RandomForestClassifier
#####
Accuracy score for RandomForestClassifier : 0.786
```




	precision	recall	f1-score	support
0	0.84	0.70	0.77	997
1	0.74	0.87	0.80	1003
accuracy			0.79	2000
macro avg	0.79	0.79	0.78	2000
weighted avg	0.79	0.79	0.78	2000

0.8985515869642827




```
##### GradientBoostingClassifier
#####
Accuracy score for GradientBoostingClassifier : 0.76
```




	precision	recall	f1-score	support
0	0.69	0.96	0.80	997
1	0.93	0.56	0.70	1003
accuracy			0.76	2000
macro avg	0.81	0.76	0.75	2000
weighted avg	0.81	0.76	0.75	2000

0.8493766443897995




```
##### BaggingClassifier
#####
Accuracy score for BaggingClassifier : 0.7595
```




	precision	recall	f1-score	support
0	0.84	0.63	0.72	997
1	0.71	0.88	0.79	1003
accuracy			0.76	2000
macro avg	0.78	0.76	0.76	2000
weighted avg	0.78	0.76	0.76	2000

0.8789239103151927




```
##### GaussianNaiveBayes
#####
Accuracy score for GaussianNaiveBayes : 0.733
```




	precision	recall	f1-score	support
0	0.81	0.61	0.70	997
1	0.69	0.85	0.76	1003
accuracy			0.73	2000
macro avg	0.75	0.73	0.73	2000
weighted avg	0.75	0.73	0.73	2000

0.7325640930768378




```
##### StochasticGradientDescent
#####
Accuracy score for StochasticGradientDescent : 0.8095
```




	precision	recall	f1-score	support
0	0.83	0.77	0.80	997
1	0.79	0.85	0.82	1003
accuracy			0.81	2000
macro avg	0.81	0.81	0.81	2000
weighted avg	0.81	0.81	0.81	2000

0.8780164021476193




```
##### KNN
#####
Accuracy score for KNN : 0.6745
```




	precision	recall	f1-score	support
0	0.75	0.52	0.61	997
1	0.63	0.83	0.72	1003
accuracy			0.67	2000
macro avg	0.69	0.67	0.67	2000
weighted avg	0.69	0.67	0.67	2000

0.806555759001831




```
##### ExtraTreesClassifier
#####
Accuracy score for ExtraTreesClassifier : 0.784
```



	precision	recall	f1-score	support
0	0.81	0.75	0.77	997
1	0.76	0.82	0.79	1003
accuracy			0.78	2000
macro avg	0.79	0.78	0.78	2000
weighted avg	0.79	0.78	0.78	2000

0.8714433429900871



CPU times: user 1h 11min 9s, sys: 4.89 s, total: 1h 11min 14s
Wall time: 1h 10min 49s

```
[44]: model_comp_df = pd.DataFrame(auc, columns = ['ROC AUC'], index = classifiers)

      model_comp_df.sort_values(by='ROC AUC', ascending=False)
```

```
[44]:
```

	ROC AUC
LogisticRegression	0.916242
SVC	0.913558
RandomForestClassifier	0.898552
BaggingClassifier	0.878924
SGDClassifier	0.878016
ExtraTreesClassifier	0.871443
GradientBoostingClassifier	0.849377

KNeighborsClassifier	0.806556
DecisionTreeClassifier	0.803740
GaussianNB	0.732564


```
[45]: from textblob import TextBlob
```

```
[46]: df_new.columns
```


```
[46]: Index(['star_rating', 'review', 'rating', 'cleaned_text'], dtype='object')
```

```
#polarity
```

```
[47]: df_new['review'] = df_new['review'].astype(str) #Make sure about the correct
      ↪ data type
      pol = lambda x: TextBlob(x).sentiment.polarity
      df_new['polarity'] = df_new['review'].apply(pol)
      import matplotlib.pyplot as plt
      import seaborn as sns
      num_bins = 50
      plt.figure(figsize=(10,6))
      n, bins, patches = plt.hist(df_new.polarity, num_bins, facecolor='blue',
      ↪ alpha=0.5)
      plt.xlabel('Polarity')
      plt.ylabel('Number of Reviews')
      plt.title('Histogram of Polarity Score')
      plt.show();
```



```
[48]: plt.figure(figsize=(10,6))
sns.boxenplot(x='rating', y='polarity', data=df_new)
plt.show();
```



#subjectivity

```
[49]: sub = lambda x: TextBlob(x).sentiment.subjectivity
df_new['subjectivity'] = df_new['review'].apply(sub)
df_new.sample(10)
```


```
[49]:
```

	star_rating	...	subjectivity
5047	5	...	0.755952
3181	4	...	1.000000
8899	5	...	0.625000
3156	1	...	0.426667
733	5	...	1.000000
9311	3	...	0.058333
8475	3	...	0.416667
4765	2	...	0.408333
9542	4	...	0.640000
8668	3	...	0.341667


[10 rows x 6 columns]

```
[50]: # Density Plot and Histogram of subjectivity
plt.figure(figsize=(10,5))
sns.distplot(df_new['subjectivity'], hist=True, kde=True,
bins=int(30), color = 'darkblue',
hist_kws={'edgecolor':'black'},
kde_kws={'linewidth': 4})
plt.xlim([-0.001,1.001])
plt.xlabel('Subjectivity', fontsize=13)
plt.ylabel('Frequency', fontsize=13)
plt.title('Distribution of Subjectivity Score', fontsize=15)
```

```
[50]: Text(0.5, 1.0, 'Distribution of Subjectivity Score')
```



```
[51]: plt.figure(figsize=(10,6))
sns.scatterplot(x='polarity', y='subjectivity', hue="rating", data=df_new)
plt.xlabel('Polarity', fontsize=13)
plt.ylabel('Subjectivity', fontsize=13)
plt.title('Polarity vs Subjectivity', fontsize=15)
plt.show();
```



[51]: